

Simon Moll / Saarland University / Saarland Informatics Campus

Introduction

The Region Vectorizer provides a single, unified API to vectorize code regions.

- RV is a generalization of the Whole-Function Vectorizer
R. Karrenberg, S. Hack, "Whole Function Vectorization" (CGO '11)

Applications

- **Outer-Loop Vectorizer** An “unroll-and-jam” vectorizer based on RV’s analysis and transformations
- **`pragma omp simd`** Emit vector code for SIMD regions right from Clang
- **Vectorizer Cost Model** How much predication? Which memory accesses vectorize well?
- **Polly** Directly vectorize loops during Polly code generation
- **PIR** Parallel region vectorizer

```
rv::VectorizationInfo vi;
// region set up
rv::Region R(xLoop);
vi.setVectorShape(xPhi,
                  VectorShape::consecutive());

// Vectorization analysis
rv::analyze(R, vi, domTree, loopInfo);

// Control conversion
rv::linearize(R, vi, domTree, loopInfo);

// Vector IR generation
rv::vectorize(R, vi, domTree);
```

| | |
|------------|--------|
| rv::Region | Region |
|------------|--------|

A **region** can be a subset of the basic blocks in a function or an entire function (`omp declare simd`).

```
#pragma omp simd
for (int x = 0; x < width; ++x) {
    for (int y = 0; y < height; ++y) {
        complex<double> c = (startX+x*step) + (startY-y*step) * I;
        complex<double> z = 0.0;

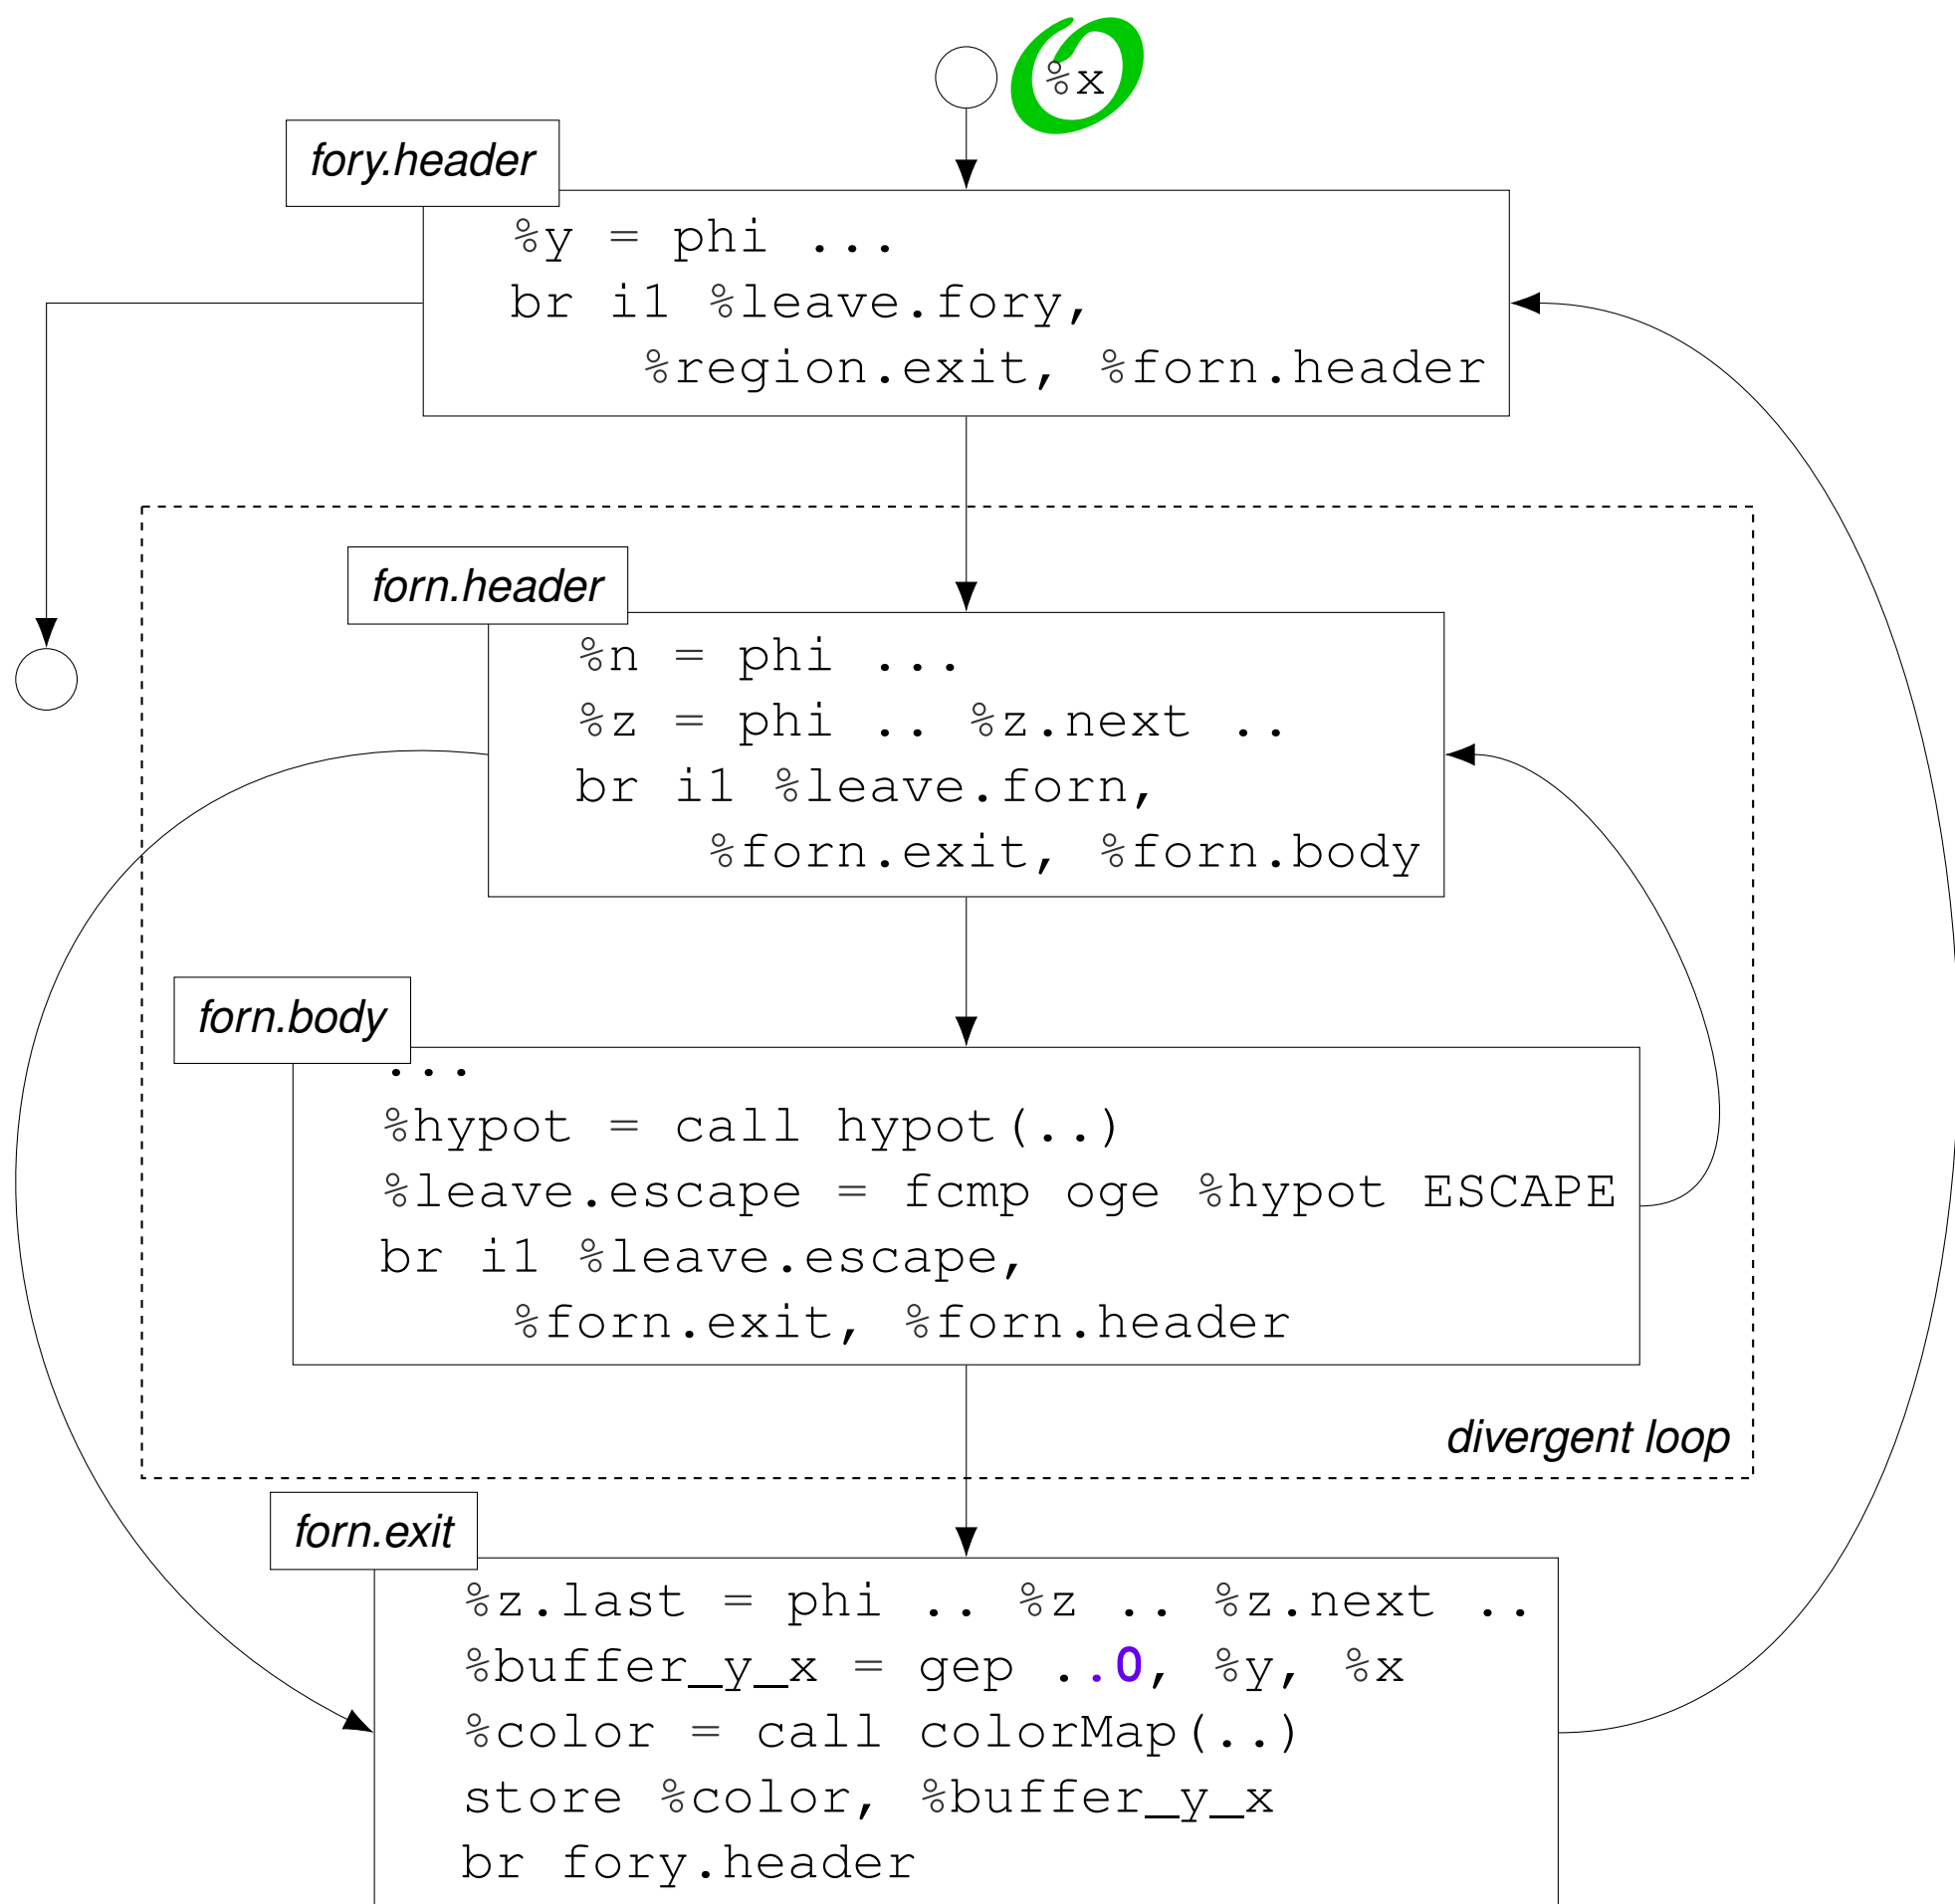
        for (int n = 0; n < MAX_ITER; ++n) {
            z = z * z + c;
            if (hypot(z.real, z.imag) >= ESCAPE)
                break;
        }
        buffer[y][x] = colorMap(z);
    }
}
```

```
#pragma omp declare simd
float min (float a, float b)
{
    if (a < b) return a; else return b;
}


↓

float min_v8 (<8 x float> a, <8 x float> b) {
    return select(a < b, a, b);
}
```

rv::analyze Vectorization Analysis




(stride, alignment) or T

| | | | | | | | | | | | | | | |
|-----|---|---|---|---|---|---|---|---|---|---|---|-----|---------------|----------------------|
| i64 |  | <table border="1"><tr><td>0</td><td>1</td><td>2</td><td>3</td></tr></table> | 0 | 1 | 2 | 3 | <table border="1"><tr><td>4</td><td>5</td><td>6</td><td>7</td></tr></table> | 4 | 5 | 6 | 7 | ... | (1, 4) | <i>(consecutive)</i> |
| 0 | 1 | 2 | 3 | | | | | | | | | | | |
| 4 | 5 | 6 | 7 | | | | | | | | | | | |
| i64 | %y, %n | <table border="1"><tr><td>3</td><td>3</td><td>3</td><td>3</td></tr></table> | 3 | 3 | 3 | 3 | <table border="1"><tr><td>4</td><td>4</td><td>4</td><td>4</td></tr></table> | 4 | 4 | 4 | 4 | ... | (0, 1) | <i>(uniform)</i> |
| 3 | 3 | 3 | 3 | | | | | | | | | | | |
| 4 | 4 | 4 | 4 | | | | | | | | | | | |
| i1 | %leave.escape | <table border="1"><tr><td>1</td><td>1</td><td>0</td><td>0</td></tr></table> | 1 | 1 | 0 | 0 | <table border="1"><tr><td>0</td><td>0</td><td>1</td><td>1</td></tr></table> | 0 | 0 | 1 | 1 | ... | ⊤ | <i>(varying)</i> |
| 1 | 1 | 0 | 0 | | | | | | | | | | | |
| 0 | 0 | 1 | 1 | | | | | | | | | | | |


Branch Divergence

Which branches cause SIMD threads to diverge?





Loop Divergence

Which loops drop off SIMD threads at different exits?



rv::linearize Control Conversion

- Optimized linearization of divergent branches and loops (\rightarrow predication)
- Preserves uniform branches and loops
- Generates **Predicated IR**
 1. All branches are uniform
 2. Blocks may be predicated



Future Work

- BOSCC (skip predicated regions if no SIMD thread is active)
J. Shin, “Introducing Control Flow into Vectorized Code” (PACT ’07)
- Multi-dimensional Analysis
C. Yount, “Vector Folding: Improving Stencil Performance via Multi-dimensional SIMD-vector Representation” (ICISS-CSS-HPCC ’15)
- Vectorization of interleaved memory accesses

- Integration with Clang / LoopVectorizer / Polly
- Reductions
 - Development available at GitHub
<https://github.com/simoll/rv>

